

MOOC 1 — Exploratory Data Analysis for Machine Learning

Original Notion Document : https://www.notion.so/MOOC-1-Exploratory-Data-Analysis-for-Machine-Learning-27cec520737c8017b4a1ec2b8ee23499?source=copy_link

Quick Table of Contents

- Module 1: EDA for Machine Learning — Overview
- Module 2: Data Retrieval and Cleaning
- Module 3: EDA and Feature Engineering
- Module 4: Estimation, Inference, and Hypothesis Testing

▼ Detailed Table of Contents

[Quick Table of Contents](#)

[Detailed Table of Contents](#)

[Module 1: Exploratory Data Analysis for Machine Learning – Study Notes](#)

- [1. Artificial Intelligence \(AI\)](#)
- [2. Machine Learning \(ML\)](#)
 - [2.1 Concept](#)
 - [2.2 Key Terms](#)
 - [2.3 Types of ML](#)
 - [2.4 Feature Engineering & Limitations](#)
- [3. Deep Learning \(DL\)](#)

4. History of AI

4.1 Foundations & Early Excitement (1950s–1990s)

4.2 First AI Winter (1960s–1970s)

4.3 1980s – Expert Systems Boom

4.4 Neural Network Progress

4.5 Second AI Winter (late 1980s–1990s)

5. Modern AI

5.1 Industry Applications

6. AI in Daily Life

6.1 Transportation

6.2 Social Media

6.3 Voice & NLP

6.4 Computer Vision & Object Detection

7. Machine Learning Workflow

7.1 Prerequisites and Tools

7.2 Typical Workflow

7.3 Basic ML Vocabulary

Module 2: Data Retrieval & Cleaning

2.1 Retrieving Data

2.2 Data Cleaning

2.2.1 Missing values

2.2.2 Duplicate data

2.2.3 Filtering and dropping

2.2.4 Outliers

Module 3: EDA and Feature Engineering — Notes with Code Examples

3.1 What is EDA?

3.2 Why is EDA useful?

3.3 EDA techniques

3.4 Sampling from DataFrames

3.5 Feature Scaling

3.6 Encoding Categorical Variables

3.6.1 Common methods

3.6.2 When to use which

3.6.3 Code with pandas

3.6.4 Code with scikit-learn

3.6.5 Practical tips

Module 4: Estimation & Inference + Hypothesis Testing — Summary Notes

4.1 Estimation


4.2 Inference

4.3 Parametric vs Non-parametric

4.4 Maximum Likelihood Estimation (MLE)

- [4.5 Common Distributions](#)
- [4.6 Frequentist vs Bayesian](#)
- [4.7 Hypothesis Testing](#)
 - [4.7.1 Concepts](#)
 - [4.7.2 Decision rule](#)
 - [4.7.3 Bayesian view](#)
 - [4.7.4 Likelihood ratio](#)
 - [4.7.5 Type I and II errors](#)
 - [4.7.6 Significance level and p-value](#)
 - [4.7.7 Coin toss example](#)
 - [4.7.8 F-statistic \(linear regression\)](#)
 - [4.7.9 Multiple testing — Bonferroni](#)
 - [4.7.10 Correlation vs Causation](#)
 - [4.7.11 Code examples \(Python\)](#)
 - [4.7.12 Frequentist vs Bayesian — Quick summary](#)
 - [4.7.13 Formulas](#)

Module 1: Exploratory Data Analysis for Machine Learning – Study Notes



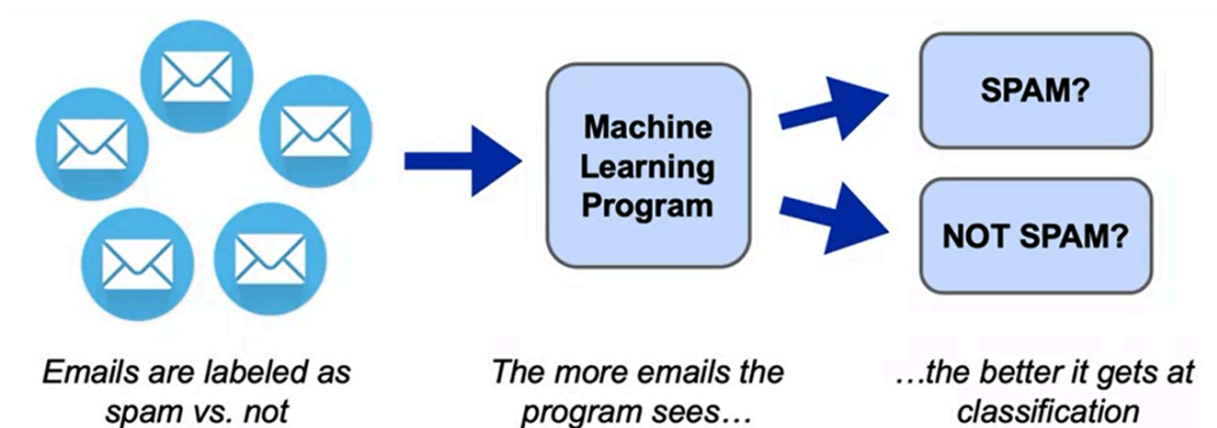
©STUDYOPEDIA All rights reserved

1. Artificial Intelligence (AI)

- **Concept:** Programs simulating intelligent behavior: sensing, reasoning, acting, adapting.
- **Details:** Not all AI learns from data; some are rule-based.
- **Examples:** Virtual assistants (Siri, Google Assistant), chatbots, rule-based systems.

- **Notes:** Subsets include ML and DL; modern AI often relies on learning from data.
- **Tip:** AI = “machines performing human-like intelligence” (Turing Test).

2. Machine Learning (ML)



2.1 Concept

- Programs learn patterns from data without explicit programming.
- More data → better learning, eventually plateaus.
- Subset of AI, includes classical ML & DL.

2.2 Key Terms

- **Feature:** Input variable used to predict target (e.g., sepal/petal length, transaction amount, pixel values).
- **Target / Label:** Output to predict (e.g., species, spam/not spam, fraud/not fraud).

2.3 Types of ML

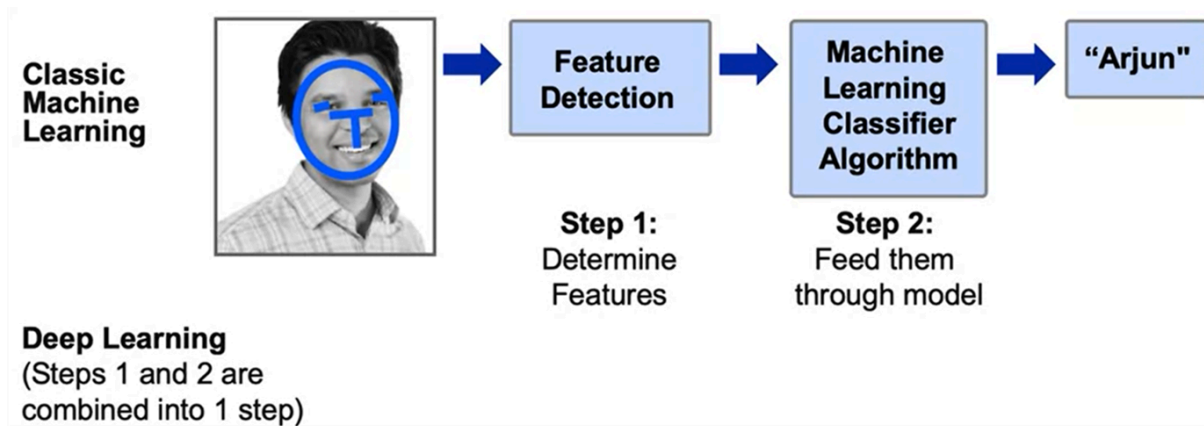
	Dataset	Goal	Example
Supervised Learning	Has a Target Column	Make Predictions	Fraud Detection
Unsupervised Learning	Does <u>not</u> have a Target Column	Find Structure in the Data	Customer Segmentation

- **Supervised Learning:** Labeled data → predict target.
 - Examples: spam detection, fraud detection, iris classification.
 - Goal: predict labels for new data.
 - Tip: Like teaching the machine with exercises that have answers.
- **Unsupervised Learning:** No labels → find hidden structure.
 - Examples: customer segmentation, clustering.
 - Goal: identify patterns or natural groupings.
 - Tip: Like exploring data to find natural clusters.

2.4 Feature Engineering & Limitations

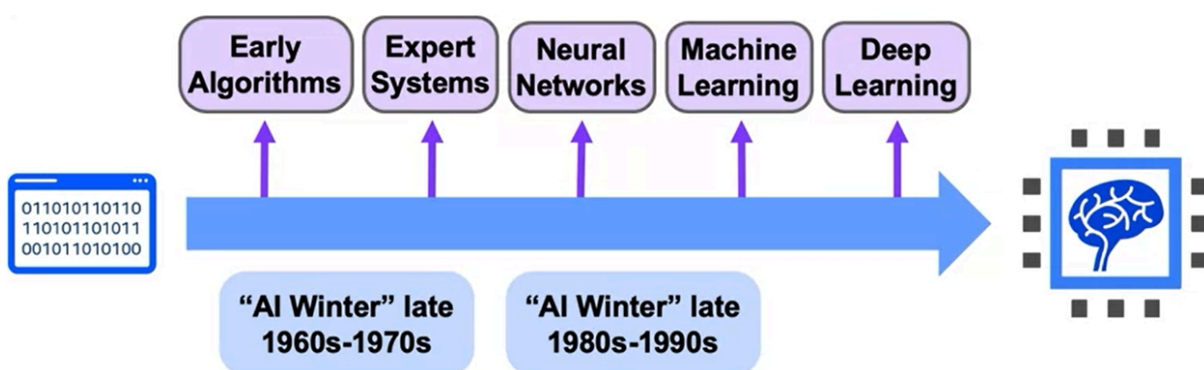
- Structured data: easier feature selection → classical ML works well.
- Unstructured data (images, audio, text): feature selection difficult → DL preferred.
- Example: Image 256×256 pixels → 65,000 features → classical ML loses spatial relationships.

3. Deep Learning (DL)



- Classical ML struggles with images:
 - Pixels as features → very high dimensionality.
 - Loses spatial relationships.
 - Requires manual feature engineering → difficult.
- **DL Advantages:**
 - Uses deep neural networks.
 - Learns features automatically.
 - Preserves spatial relationships.
- **Layer Workflow:** edges → shapes → high-level objects → prediction.

4. History of AI



4.1 Foundations & Early Excitement (1950s–1990s)

- **1950:** Alan Turing → Turing Test (can machines imitate humans?).
- **1956:** Dartmouth Conference → AI coined; goal: simulate intelligent behavior.

- **1957:** Rosenblatt → Perceptron → first neural network, learned from data.
- **1959:** Arthur Samuel → Checkers program → learns from past moves; popularized Machine Learning.

4.2 First AI Winter (1960s–1970s)

- Machine translation fails (Russian↔English) → low ROI → US government scrutiny.
- **1969:** Minsky → limitations of perceptron.
- **1973:** Lighthill report → research fell short → major funding cuts.

4.3 1980s – Expert Systems Boom

- Rule-based systems mimicking human experts.
- Ran on mainframes using LISP.
- Applied in business for decision-making → AI practically useful.

4.4 Neural Network Progress

- Geoffrey Hinton → Backpropagation algorithm → multi-layer networks can learn.
- Theoretical breakthroughs → potential for deep learning.

4.5 Second AI Winter (late 1980s–1990s)

- Expert systems unable to learn; brittle with unusual inputs.
- Neural networks couldn't scale; backpropagation had problems with large datasets.
- Investment decline as AI hype subsided.

5. Modern AI

- **Growth Areas:** Computer Vision (self-driving cars, medical imaging), NLP (translation, sentiment, text generation).
- **Why This Era is Different:** Larger datasets (ImageNet), faster & cheaper computing, neural networks & deep learning → practical applications.

5.1 Industry Applications

- **Healthcare:** Diagnosis, drug discovery, sensory aids.
- **Industrial:** Automation, predictive maintenance, agriculture.
- **Finance:** Trading, fraud detection, risk management.
- **Energy:** Smart grids, resource extraction, conservation.
- **Government:** Defense, citizen services, smart cities.
- **Transportation:** Autonomous vehicles, logistics, drones.
- **Education & Entertainment:** Personalized learning, gaming, media.

6. AI in Daily Life

6.1 Transportation

- Route optimization: Google Maps, Waze predict traffic, fastest routes.
- Ride-sharing: Uber, Lyft adjust pricing in real-time (supply & demand).

6.2 Social Media

- Content personalization: posts, groups, targeted ads.
- Sentiment analysis: reviews or content mood.
- Image recognition: face/object tagging and sharing.

6.3 Voice & NLP

- Virtual assistants: Siri, Alexa understand commands.
- NLP: translation, sentiment detection, task automation.

6.4 Computer Vision & Object Detection

- Self-driving cars: live detection of objects, pedestrians, road conditions.
- Security: detect abandoned baggage in real time.
- Image classification: DL can outperform humans.

7. Machine Learning Workflow


7.1 Prerequisites and Tools

- Python libraries: NumPy, Pandas, Matplotlib, Seaborn, Scikit-Learn, TensorFlow, Keras.

- Environment: Jupyter Notebook / iPython.
- Math background: basic statistics, probability, Bayes' rule, linear algebra.

7.2 Typical Workflow

1. **Problem Statement:** Define the problem (e.g., classify dog breeds).
2. **Data Collection:** Gather labeled data (variety of angles, lighting, etc.).
3. **Data Exploration & Preprocessing:**
 - Clean data, handle missing values.
 - Analyze distributions, correlations, heatmaps.
 - Convert data into suitable formats (e.g., pixel arrays).
4. **Modeling:**
 - Build baseline models.
 - Train on training data.
5. **Validation:**
 - Evaluate using holdout/test set.
 - Measure accuracy or other metrics.
6. **Decision-making & Deployment:**
 - Communicate results, deploy model.



7.3 Basic ML Vocabulary

- **Target variable:** Value to predict (e.g., iris species).
- **Features / Explanatory variables:** Columns used to predict target (e.g., petal/sepal length & width).
- **Example / Observation:** Single row in dataset.
- **Label:** Value of target variable for one example (e.g., "versicolor").

Module 2: Data Retrieval & Cleaning

2.1 Retrieving Data

Key points

- Common data sources:
 - SQL databases: MySQL, PostgreSQL, SQLite
 - NoSQL databases: MongoDB, Redis, Cassandra
 - APIs: REST, JSON-based services
 - Cloud data sources: AWS, GCP, Azure
- Common formats:
 - CSV or TSV: tabular data with commas or tabs
 - JSON: key-value format, similar to a Python dict

Overview table

Type	Format	Python library	Typical use case
CSV	Comma separated tabular data	<code>pandas</code> → <code>'pd.read' _csv()</code>	General tabular datasets
JSON	Key-value objects and arrays	<code>pandas</code> → <code>'pd.read' _json()</code>	APIs and NoSQL exports
SQL	Relational tables queried with SQL	<code>sqlite3</code> , <code>pandas</code> → <code>'pd.read' _sql()</code>	MySQL, PostgreSQL, SQLite
NoSQL	Document or key-value stores	<code>pymongo</code> (MongoDB), drivers per DB	MongoDB, Redis, Cassandra

2.2 Data Cleaning

Key points

- Garbage in = Garbage out. Dirty data leads to unreliable models.
- Main issues:
 - Missing values
 - Duplicate data
 - Outliers
 - Unnecessary or irrelevant data

2.2.1 Missing values

Options

- Remove: drop rows with NaN, but you lose data
- Impute: replace with mean or median, adds uncertainty
- Mask: assign a category for missing values

Syntax

```
# Remove
df.dropna()

# Impute with mean
df.fillna(df.mean())

# Impute with median
df.fillna(df.median())

# Mask as a category
df.fillna("Missing")
```

2.2.2 Duplicate data

Key points

- Duplicates bias metrics and can leak across train and test splits.
- Causes include rerun pipelines and multiple manual entries.
- Solutions: drop all duplicates or enforce uniqueness by specific column sets.

Syntax

```
# Drop all duplicate rows
df.drop_duplicates()

# Enforce uniqueness by multiple columns
df.drop_duplicates(subset=["FirstName", "LastName"])

# Enforce uniqueness by a single column
df.drop_duplicates(subset=["brand"])
```

2.2.3 Filtering and dropping

Key points

- Remove irrelevant data and keep only what's needed.
- Filter by name, regex, or substring.

Syntax

```
# Select columns by explicit names
df.filter(items=['one', 'three'])

# Select columns ending with 'e'
df.filter(regex='e$', axis=1)

# Select rows containing substring 'bbi' in the index
df.filter(like='bbi', axis=0)

# Drop columns B and C
df.drop(['B', 'C'], axis=1)
```

2.2.4 Outliers

Key points

- Outliers are far from most observations and strongly affect many models.
- Detection methods: plots, IQR, Z-score, residuals.
- Handling: remove, impute, transform, or use robust models.

Syntax

```
# IQR method
Q1 = df['col'].quantile(0.25)
Q3 = df['col'].quantile(0.75)
IQR = Q3 - Q1
outliers = df[(df['col'] < (Q1 - 1.5 * IQR)) | (df['col'] > (Q3 + 1.5 * IQR))]
```

```
# Z-score method
import numpy as np
```

```
z_scores = np.abs((df['col'] - df['col'].mean()) / df['col'].std())
outliers = df[z_scores > 3]
```

Module 3: EDA and Feature Engineering — Notes with Code Examples

3.1 What is EDA?

- Exploratory Data Analysis summarizes the main characteristics of a dataset using descriptive statistics and visualizations.
- Goal: understand structure, quality, and patterns before modeling.

```
import pandas as pd

# Example: load a CSV (replace with your dataset)
# csv_url = "https://.../your_data.csv"
# df =
```

3.2 Why is EDA useful?

- Get an initial feel for the data
- Detect cleaning and preprocessing needs
- Spot errors, missing values, inconsistencies
- Reveal patterns, trends, and relationships

```
# Structural overview
print(
```

3.3 EDA techniques

- Descriptive stats: mean, median, min, max, variance, correlations

```
# Average Age by Sex (use columns that exist in your dataset)
print(df.groupby("Sex")["Age"].mean())
```

```
# Correlation matrix (numeric only)
print(df.corr(numeric_only=True))
```

- Visualization: histograms, scatter plots, box plots, pair plots

```
import matplotlib.pyplot as plt
import seaborn as sns

# Histogram of Age
_ = df["Age"].hist(bins=30)
plt.xlabel("Age")
plt.ylabel("Frequency")
plt.title("Distribution of Age")
```

3.4 Sampling from DataFrames

- Reasons: speed up experimentation, balance rare outcomes, quick tests on subsets

```
# Random sample of 100 rows
sample_df = df.sample(100, random_state=42)
print(sample_df.head())
```

3.5 Feature Scaling

Why scaling?

- Different features can have very different ranges. For example, $\text{Age} \in [0, 100]$ vs $\text{Income} \in [0, 100000]$.
- Distance- and gradient-based algorithms such as KNN, SVM, Logistic Regression, and Neural Networks are scale-sensitive. Unscaled large-range features can dominate the optimization.
- Scaling puts features on comparable ranges to stabilize training and improve model performance.

Scaling methods

1) Standard scaling (Z-score normalization)

- Formula: $z = \frac{x - \mu}{\sigma}$

- Effect: centers features to mean 0 and standard deviation 1
- Notes: often a good default for many linear models and neural networks

```
from sklearn.preprocessing import StandardScaler
X_std = StandardScaler().fit_transform(X)
```

2) Min–Max scaling

- Formula: $x' = \frac{x - x_{\min}}{x_{\max} - x_{\min}}$ √x Phương trình mới
- Effect: maps features to [0, 1]
- Notes: useful when models expect bounded inputs; sensitive to outliers

```
from sklearn.preprocessing import MinMaxScaler
X_mm = MinMaxScaler().fit_transform(X)
```

3) Robust scaling

- Formula (per feature): $x' = \frac{x - \mathrm{median}(x)}{\mathrm{IQR}(x)}$
\$, where $\mathrm{IQR} = Q_{75} - Q_{25}$
- Effect: centers by the median and scales by IQR
- Notes: less sensitive to outliers than Standard or Min–Max


```
from sklearn.preprocessing import RobustScaler
X_robust = RobustScaler().fit_transform(X)
```

When to use which?

- If features have outliers: prefer RobustScaler
- If a bounded [0, 1] range helps or a model expects it: MinMaxScaler
- Otherwise, as a default starting point: StandardScaler

Practical tips

- Fit scalers on the training set only, then transform validation and test sets with the fitted scaler to avoid leakage
- Persist the fitted scaler for production to apply the exact same transformation

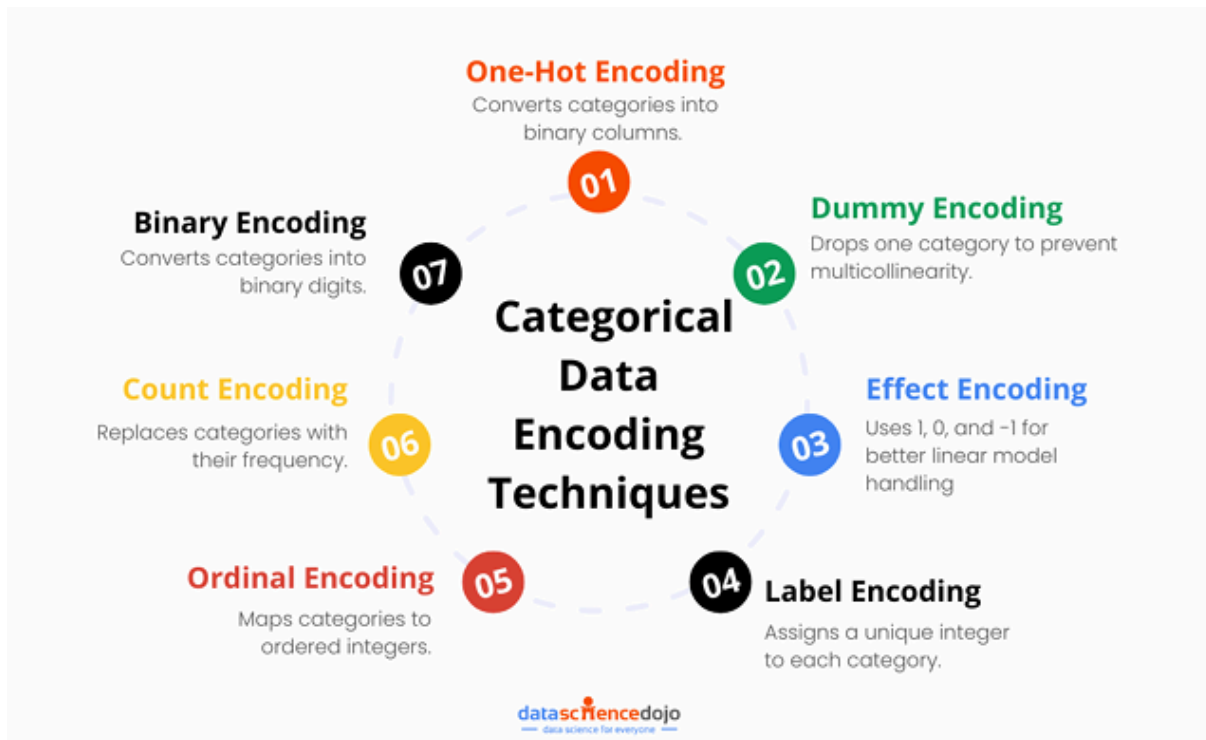


3.6 Encoding Categorical Variables

- Goal: transform categorical variables into numeric representations that models can learn from.

3.6.1 Common methods

- Binary encoding: two states \rightarrow 0/1. Useful for flags and booleans.
- One-hot encoding: each category becomes a 0/1 column. Avoids imposing fake order.
- Dummy encoding: like one-hot but drop one reference column to avoid multicollinearity with an intercept.
- Ordinal encoding: map a true order Low < Medium < High \rightarrow 1/2/3.
- Label encoding: integers 0..K-1 for categories. Prefer for tree-based models or truly ordered categories.



3.6.2 When to use which

- No natural order and using linear, KNN, SVM, or neural nets → One-hot or Dummy.
- Natural order present → Ordinal.
- High-cardinality features → consider Target encoding, Hashing encoding, or grouping rare levels.
- Tree models (GBDT/Random Forest) can work with Label/Ordinal, but beware introducing artificial order.

3.6.3 Code with pandas

```
import pandas as pd

# One-hot (keep all columns)
df_oh = pd.get_dummies(df, columns=["category"], dtype=int)

# Dummy (drop reference column to avoid multicollinearity)
df_dummy = pd.get_dummies(df, columns=["category"], drop_first=True, dtype=int)
```

3.6.4 Code with scikit-learn

```
from sklearn.preprocessing import OneHotEncoder, OrdinalEncoder

# One-hot encoder (robust to unseen categories at inference)
ohe = OneHotEncoder(handle_unknown="ignore", sparse_output=False)
X_oh =
```

3.6.5 Practical tips

- Build a pipeline: fit encoders on the training set, then transform validation and test to avoid leakage.
- Persist fitted encoders with the model for production.
- Combine rare categories into "Other" to reduce one-hot width and stabilize estimates.
- Guard against unseen categories at inference: use `handle_unknown="ignore"` for `OneHotEncoder` or map to "Other" before encoding.
- Purpose: correct skewness, reduce outlier impact, linearize relationships, and help optimization converge
- Log transformation
 - Use for right-skewed variables such as income; use `log1p` when zeros are present

```
import numpy as np
import matplotlib.pyplot as plt

x = np.random.exponential(scale=2, size=1000)
x_log = np.log1p(x)

plt.figure(figsize=(12,4))
plt.subplot(1,2,1); plt.hist(x, bins=30, color="skyblue"); plt.title("Original (Right-Skewed)")
plt.subplot(1,2,2); plt.hist(x_log, bins=30, color="orange"); plt.title("Log-Transformed")
```

- Polynomial features
 - Use when relationships are non-linear; beware of overfitting at high degrees

```
from sklearn.preprocessing import PolynomialFeatures
import numpy as np
import pandas as pd

X = np.arange(6).reshape(-1, 1)
poly = PolynomialFeatures(degree=2, include_bias=False)
X_poly =
```

- Feature engineering improves model accuracy by creating and transforming variables
- Common: logs and polynomials for skewness and nonlinearity. Encode categoricals. Normalize for comparability
- Principles: understand variable types and model assumptions. Apply consistently to train and test

Module 4: Estimation & Inference + Hypothesis Testing — Summary Notes

4.1 Estimation

- **Sample mean:** $\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$
- **Sample variance:** $S^2 = \frac{\sum_{i=1}^n (X_i - \bar{X})^2}{n-1}$
- **Sample standard deviation:** $S = \sqrt{\frac{\sum_{i=1}^n (X_i - \bar{X})^2}{n-1}}$

4.2 Inference

- **Standard Error (SE):** $SE = \frac{S}{\sqrt{n}}$
- **t-CI for mean:** $\bar{X} \pm t_{\alpha/2, n-1} \cdot SE$

4.3 Parametric vs Non-parametric

- **Parametric:** assumes a distribution with a finite number of parameters.
Example: Normal (μ, σ^2)
- **Non-parametric:** fewer assumptions, not distribution-based. Examples: histogram, empirical CDF, KDE

4.4 Maximum Likelihood Estimation (MLE)

- **Likelihood:** $L(\theta \mid x) = \prod_{i=1}^n f(x_i \mid \theta)$
- **Log-likelihood:** $\ell(\theta) = \sum_{i=1}^n \ln f(x_i \mid \theta)$
- **MLE estimator:** $\hat{\theta}_{\text{MLE}} = \arg \max_{\theta} \ell(\theta)$

4.5 Common Distributions

- Uniform on $[a, b]$ with $a < b$ and density:

$$f(x) = \frac{1}{b-a}, \quad a \leq x \leq b$$

- Normal with mean μ và variance σ^2 :

$$f(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$$

- Exponential với tốc độ $\lambda > 0$:

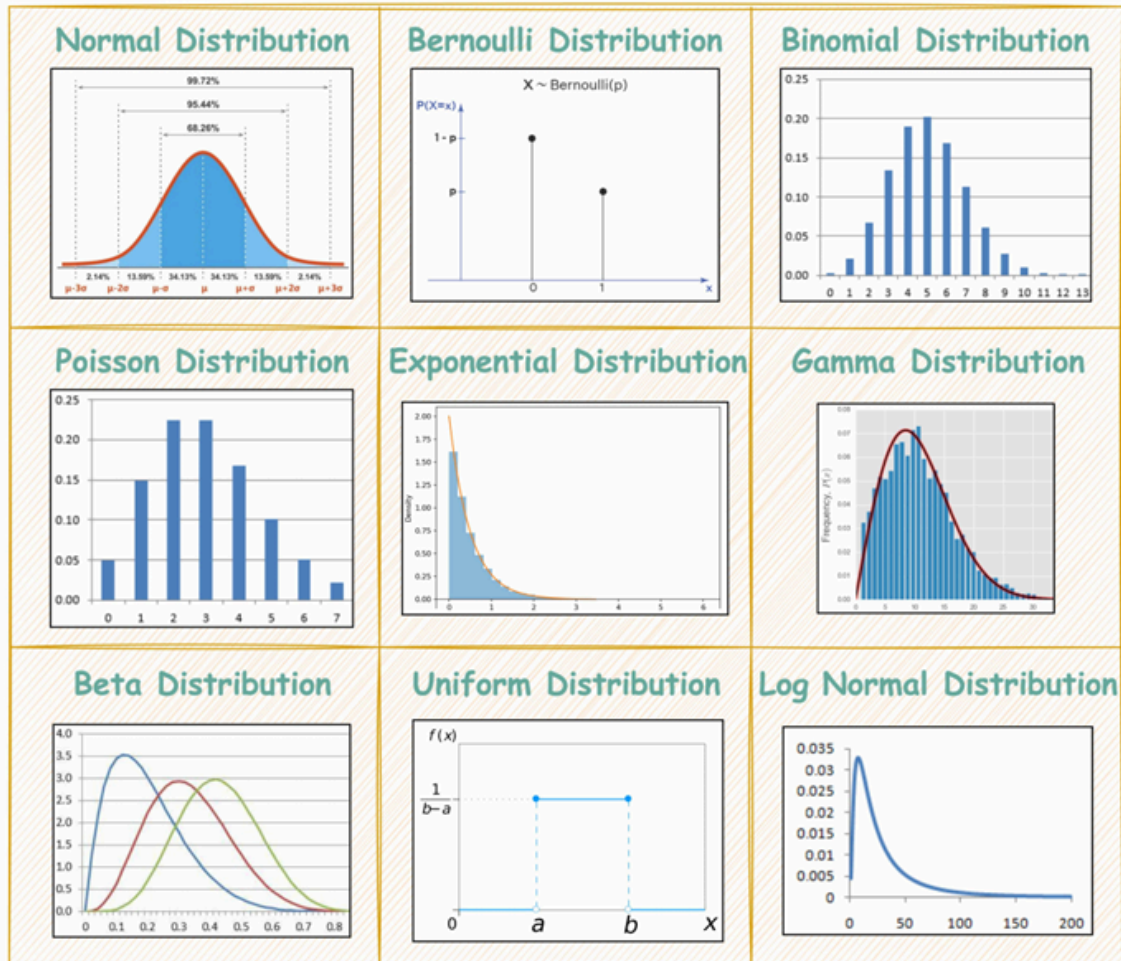
$$f(x) = \lambda e^{-\lambda x}, \quad x \geq 0$$

- Poisson với tốc độ $\lambda > 0$:

$$\mathbb{P}(X = k) = \frac{\lambda^k e^{-\lambda}}{k!}, \quad k = 0, 1, 2, \dots$$



Most Important Distributions in Data Science



4.6 Frequentist vs Bayesian

- **Frequentist:** Probability = long-run frequency under repeated experiments
- **Bayesian:** Parameters are random variables. Bayes: $P(\theta | D) = \frac{P(D | \theta) P(\theta)}{P(D)}$

4.7 Hypothesis Testing

4.7.1 Concepts

- H_0 (null): baseline hypothesis, no effect
- H_1 (alternative): opposing hypothesis, effect exists

4.7.2 Decision rule

- From a test statistic define the **rejection region** for H_0 and the **acceptance region**

4.7.3 Bayesian view

- $P(H_i | D) = \frac{P(D | H_i) P(H_i)}{P(D)}$

4.7.4 Likelihood ratio

- $\Lambda(x) = \frac{P(\text{data} | H_1)}{P(\text{data} | H_0)}$

4.7.5 Type I and II errors

- **Type I (α):** reject H_0 when H_0 is true
- **Type II (β):** fail to reject H_0 when H_0 is false
- **Power:** $1 - \beta$

4.7.6 Significance level and p-value

- Choose α commonly 0.05 or 0.01
- Rule: reject H_0 if $p < \alpha$

4.7.7 Coin toss example

- $\$H_0\$$: fair coin $\Rightarrow X \sim \text{Binomial}(n = 10, p = 0.5)$
- Observe 3 heads. $p\text{-value} = P(X \leq 3) \approx 0.17 > 0.05 \Rightarrow$ **do not reject** H_0

4.7.8 F-statistic (linear regression)

- $\$H_0\$$: all coefficients $\beta = 0$ (cụ thể: $\$beta_1=beta_2=\cdots=beta_p=0\$$).
Reject if F-test p-value is small enough

4.7.9 Multiple testing — Bonferroni

- Adjust significance level: $\alpha' = \frac{\alpha}{m}$

4.7.10 Correlation vs Causation

- Correlation \neq causation

4.7.11 Code examples (Python)

- CI for the mean (use t-distribution when σ is unknown):

```
import numpy as np
from scipy import stats

# Example sample
x = np.array([12.1, 11.7, 12.4, 11.9, 12.3, 12.0, 11.8, 12.2])

n = len(x)
mean = x.mean()
s = x.std(ddof=1)
alpha = 0.05
se = s / np.sqrt(n)

# t critical with df = n-1
t_crit = stats.t.ppf(1 - alpha/2, df=n-1)
ci = (mean - t_crit*se, mean + t_crit*se)
mean, s, ci
```

- One-sample test ($H_0: \mu = \mu_0$) using t-test:

```
mu0 = 12.0
# two-sided t-test
res = stats.ttest_1samp(x, popmean=mu0)
res.statistic, res.pvalue
```

- Coin p-value example ($n=10$, $k=3$, $H_0: p=0.5$, two-sided or one-sided):

```
from scipy.stats import binom
n, k, p0 = 10, 3, 0.5
# one-sided p-value:  $P(X \leq k)$ 
p_left = binom.cdf(k, n, p0)
# one-sided p-value:  $P(X \geq k)$ 
p_right = 1 - binom.cdf(k-1, n, p0)
# two-sided p-value (simple approach: double the more "extreme" side nea
```



```

r the mean)
p_two_sided = 2 * min(p_left, p_right)
p_left, p_right, p_two_sided

```

4.7.12 Frequentist vs Bayesian — Quick summary

Aspect	Frequentist	Bayesian
Meaning of probability	Long-run frequency of events	Degree of belief, prior-informed
Parameters	Fixed but unknown	Random variables with a prior distribution
Inference	Based on sampling distribution of statistics	Prior + likelihood → posterior
Intervals	CI: procedure yields 1-α coverage in repeated samples	Credible interval: probability the parameter lies in the interval
Decision	p-values, significance level α, hypothesis tests	Maximize or integrate over posterior, Bayes factor

4.7.13 Formulas

$$SE = \frac{S}{\sqrt{n}} \quad CI_{\mu} :$$

$$\bar{X} \pm t_{\alpha/2, n-1} \cdot \frac{S}{\sqrt{n}}$$

$$\Lambda(x) = \frac{\mathcal{L}(H_1)}{\mathcal{L}(H_0)} \quad \text{Bonferroni: } \alpha' = \frac{\alpha}{m}$$

- $X \rightarrow Y$
- $Y \rightarrow X$
- Confounders affect both
- Spurious correlation