

Modules 6: Model Interpretability & Unbalanced Classes

A consolidated document on machine learning model interpretability and techniques for handling imbalanced data.

1. Concept and importance of Model Interpretability

Concept

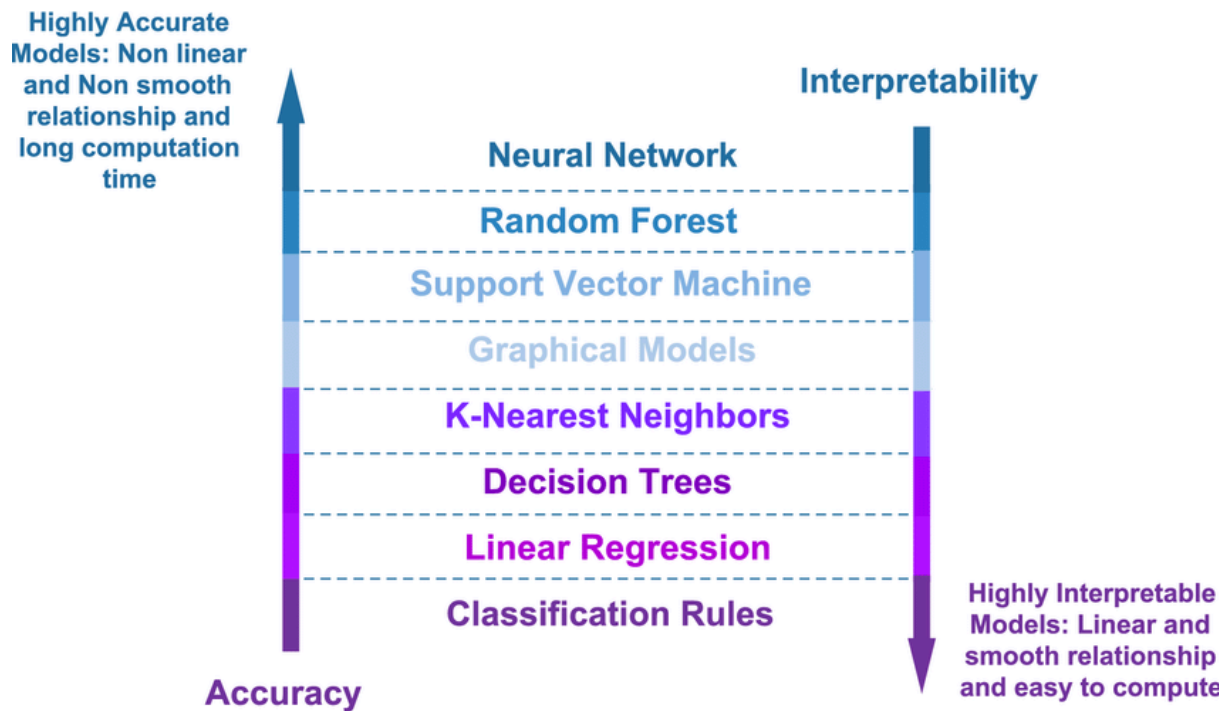
Model interpretability is the degree to which humans can understand **how an ML model makes predictions**.

It helps build **trust**, enables **debugging**, and supports **decision-making** based on the model.

Why it matters

- Understand why the model makes mistakes
 - Increase transparency in sensitive domains (healthcare, finance, HR)
 - Support model optimization and reduce bias
-

2. Model types by interpretability



Model type	Interpretability	Examples
Self-interpretable models	Easy to understand, clear structure	Linear Regression, Decision Tree, KNN
Non-self-interpretable models (Black-box)	Hard to understand, complex	Random Forest, Gradient Boosting, SVM, Deep Neural Network

◆ Self-interpretable Models

1. Linear models

Formula:

$$y = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$$

Explanation: Each coefficient w_i reflects the influence of feature x_i on the outcome.

Example: If $w_{\text{age}} = 0.5$, increasing age by 1 unit increases the log-odds by 0.5.

2. Tree models

Understandable via **if-else condition branches**.

Example:

```
if income > 5000:
    predict = "Buy"
```

```
else:  
    predict = "Not Buy"
```

3. K-nearest neighbors (KNN)

Explainable by **nearby data points**.

Example: If 4 out of 5 neighbors are "Yes", the new point is predicted as "Yes".

3. Model interpretation methods

Two main families

Family	Applies to	Goal
Intrinsic methods	Self-explaining models (Linear, Tree, KNN)	Understand directly from model structure
Post-hoc methods	Black-box models	Explain after training

4. Model-agnostic explanation methods



(a) Permutation Feature Importance (PFI)

Measures feature importance by **shuffling a feature's values** and observing how much performance drops.

General formula:

$$\text{Importance}(X_i) = \text{Score}_{\text{original}} - \text{Score}_{\text{shuffled}(X_i)}$$

Where:

- Score_original: baseline score
- Score_shuffled(X_i): score after shuffling feature X_i

Example:

Feature	Accuracy (original)	Accuracy (shuffled)	Importance
Age	0.90	0.75	0.15
Salary	0.90	0.80	0.10

→ "Age" is more important than "Salary".

(b) Partial Dependence Plot (PDP)

Describes the **relationship between one or several features and the model output** while averaging over other features.

Formula:

$$\text{PDP}(x_j) = \frac{1}{n} \sum_{i=1}^n f(x_j, x_{i,-j})$$

Where:

- $f(x_j, x_{i,-j})$: model output when fixing x_j and keeping other features as in sample i
- n : number of samples

Example:

The PDP of experience shows: as experience increases from 0 to 10 years, the probability of "hired" increases then saturates.

Ví dụ minh họa PDP (với scikit-learn):

Giả sử ta có bộ dữ liệu phân loại nhị phân với đặc trưng `experience` và các đặc trưng khác. Ta huấn luyện một mô hình Gradient Boosting rồi vẽ PDP cho `experience`.

```
import numpy as np
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.inspection import PartialDependenceDisplay
import matplotlib.pyplot as plt

# 1) Tạo dữ liệu mẫu
X, y = make_classification(
    n_samples=2000,
    n_features=5,
    n_informative=3,
    n_redundant=0,
    random_state=42,
)
feature_names = ["experience", "salary", "age", "overtime", "dept_size"]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, stratify=y, random_state=42
)
```

```
# 2) Huấn luyện mô hình
clf = GradientBoostingClassifier(random_state=42)
clf.fit(X_train, y_train)

# 3) Vẽ PDP cho đặc trưng 'experience' (giả sử là cột 0)
fig, ax = plt.subplots(figsize=(6, 4))
PartialDependenceDisplay.from_estimator(
    clf, X_test, features=[0], feature_names=feature_names, ax=ax
)
ax.set_title("PDP cho 'experience'")
plt.tight_layout()
plt.show()
```

Diễn giải: Đường cong PDP cho biết xu hướng trung bình của xác suất dự đoán theo **experience** khi giữ các đặc trưng còn lại như ở từng mẫu trong tập kiểm thử. Nếu đường cong tăng dần rồi bão hòa, điều đó khớp với mô tả ở trên.

🔮 (c) Surrogate models

Use a **simple model (Linear/Tree)** to mimic and explain a **complex black-box model**.

Two types:

1. **Global surrogate:** mimic the entire model
2. **Local surrogate:** mimic locally around one data point

Examples:

- Use a Decision Tree to explain XGBoost predictions globally
- Use LIME to explain a specific instance locally

💬 (d) LIME (Local Interpretable Model-Agnostic Explanations)

Explains **individual predictions** by fitting a **small linear model** around the instance of interest.

Generates perturbed samples near the instance, observes output changes, and trains a simple model locally.

Example:



Prediction "employee will quit" → LIME highlights "Overtime = Yes" and "YearsAtCompany < 2" as top contributors.

5. Models for imbalanced classes

The problem

When class counts are skewed, models tend to favor the majority class.

Remedies

Method	Description
Downsampling	Reduce the number of majority samples
Upsampling	Duplicate or synthesize minority samples
Combine both	Balance by adjusting both sides

Example:

- "Non-fraud": 900 samples
- "Fraud": 100 samples

→ Either upsample "Fraud" to 900 or downsample "Non-fraud" to 100.

6. Techniques for handling imbalanced data

2.1. Downsampling (reduce majority class)

Keep the minority class size, **randomly drop** majority samples.

Pros

- Simple to implement
- Reduces majority bias

Cons

- Information loss (discarding real data)

Python example:

```
from sklearn.utils import resample
import numpy as np

X_majority = X[y == 0]
X_minority = X[y == 1]
```

```

X_majority_downsampled = resample(
    X_majority,
    replace=False,
    n_samples=len(X_minority),
    random_state=42,
)

X_balanced = np.vstack((X_majority_downsampled, X_minority))
y_balanced = np.hstack(([0]*len(X_minority), [1]*len(X_minority)))

```

2.2. Upsampling (increase minority class)

Duplicate or generate new samples for the minority class.

✓ Pros

- Preserve all majority data
- Improve recall for the minority class

✗ Cons

- May cause overfitting

Python example:

```

from sklearn.utils import resample
import numpy as np

X_minority_upsampled = resample(
    X_minority,
    replace=True,
    n_samples=len(X_majority),
    random_state=42,
)

X_balanced = np.vstack((X_majority, X_minority_upsampled))
y_balanced = np.hstack(([0]*len(X_majority), [1]*len(X_majority)))

```

2.3. Hybrid methods (SMOTE, ADASYN, Tomek Links, NearMiss)

◆ SMOTE (Synthetic Minority Oversampling Technique)

Create synthetic points by **interpolating between a minority point and one of its nearest neighbors (KNN)**.

Variants:

- **SMOTE-Regular:** pick K neighbors at random
- **Borderline-SMOTE:** synthesize near the decision boundary
- **SMOTE-SVM:** use support vectors to synthesize

Formula:

$$\mathbf{x}_{\text{new}} = \mathbf{x}_i + \lambda (\mathbf{x}_{z_i} - \mathbf{x}_i), \quad \lambda \in [0, 1]$$

Code example:

```
from imblearn.over_sampling import SMOTE

smote = SMOTE(random_state=42)
X_res, y_res = smote.fit_resample(X, y)
```

◆ ADASYN (Adaptive Synthetic Sampling)

- Generate more synthetic samples **where classification is harder** (more nearby opposite-class points)
- Focus on ambiguous regions

Code example:

```
from imblearn.over_sampling import ADASYN

adasyn = ADASYN(random_state=42)
X_res, y_res = adasyn.fit_resample(X, y)
```

◆ Advanced undersampling

NearMiss:

- Keep majority samples **closest to the minority** (Euclidean distance)
- Three versions: NearMiss-1, NearMiss-2, NearMiss-3

Tomek Links:

- Remove nearest-opposite-class pairs to **clean the boundary**

Code example:


```

from imblearn.under_sampling import NearMiss, TomekLinks

nm = NearMiss(version=1)
X_res, y_res = nm.fit_resample(X, y)

tl = TomekLinks()
X_res2, y_res2 = tl.fit_resample(X, y)

```

◆ Combinations (SMOTE + Tomek Links / SMOTE + ENN)

- **SMOTE + Tomek Links:** upsample then clean boundary
- **SMOTE + ENN (Edited Nearest Neighbors):** remove noisy points after synthesis

```

from imblearn.combine import SMOTETomek, SMOTEENN

smote_tomek = SMOTETomek(random_state=42)
X_res, y_res = smote_tomek.fit_resample(X, y)

smote_enn = SMOTEENN(random_state=42)
X_res2, y_res2 = smote_enn.fit_resample(X, y)

```

◆ Balanced Bagging ("Blagging")

- Each bag in bagging is constructed so **classes are balanced**
- Improves stability and reduces bias in ensembles

Example:

```

from imblearn.ensemble import BalancedBaggingClassifier
from sklearn.tree import DecisionTreeClassifier

bbc = BalancedBaggingClassifier(
    estimator=DecisionTreeClassifier(),
    n_estimators=10,
    random_state=42,
)

```

2.4. Class weighting

Many models support **class weights**:

$$w_c = \frac{N}{N_c}$$

Where:

- N: total number of samples
- N_c: number of samples in class *c*

Example with Logistic Regression:

```
from sklearn.linear_model import LogisticRegression

model = LogisticRegression(class_weight='balanced', max_iter=1000)
```

2.5. Stratified sampling

Use `stratify` to keep **class ratios consistent** between train and test sets.

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X, y, stratify=y, random_state=42
)
```

7. Evaluation for imbalanced data



Avoid plain Accuracy as it can be misleading.

Prefer:

- **Precision, Recall, F1-score**
- **AUC (Area Under the ROC Curve)**
- **Cohen's Kappa**

Code example:

```
from sklearn.metrics import classification_report, roc_auc_score

y_pred = model.predict(X_test)
```

```
print(classification_report(y_test, y_pred))
print("AUC:", roc_auc_score(y_test, model.predict_proba(X_test)[: , 1]))
```



Summary

Technique	Goal	Pros	Cons	Key code
Downsampling	Reduce majority samples	Simple	Information loss	<code>resample(..., replace=False)</code>
Upsampling	Increase minority samples	Keep all data	Overfitting	<code>resample(..., replace=True)</code>
SMOTE	Synthesize data	Effective	May add noise	<code>SMOTE()</code>
ADASYN	Adaptive synthesis	Good near boundary	More complex	<code>ADASYN()</code>
NearMiss	Filter majority	Sharper boundary	Possible under-representation	<code>NearMiss()</code>
Tomek Links	Clean boundary	Less noise	No data synthesis	<code>TomekLinks()</code>
Blagging	Balanced ensemble	Stable	Resource intensive	<code>BalancedBaggingClassifier()</code>

Key takeaways

Topic	Essence
Interpretability goals	Understand models, build trust, detect errors
Interpretable models	Linear, Tree, KNN
Post-hoc methods	PFI, PDP, Surrogate, LIME
Imbalanced data	Use upsampling or downsampling to improve performance