

Graded Quiz: Module 1 - Modern AI and its Applications

Graded Assignment • 30 min

← Back

Due

Sep 24, 11:59 PM PDT

Your grade: 100%

Your latest: 100%

•

Your highest: 100%

To pass you need at least 70%. We keep your highest score.

Next item →

1. What is the goal of supervised learning?

Predict the features.

Predict the labels.

Find an underlying structure of the dataset without any labels.

Find the target.

Correct

The goal for supervised learning is to be able to predict the label.

1 / 1 point

2. What is deep learning?

Deep learning is machine learning that involves deep neural networks.

Deep learning is another name for artificial intelligence.

Deep learning includes artificial intelligence and machine learning.

None of the above are correct.

Correct

Deep learning is machine learning that involves using very complicated models called deep neural networks. Deep learning is a subset of machine learning.

1 / 1 point

3. When is a standard machine learning algorithm usually a better choice than using deep learning to get the job done?

When working with small data sets.

When the data is steady over time.

When working with large data sets.

None of the above are correct.

Correct

A standard machine learning algorithm is a better choice when you are working with smaller datasets, and if the data is changing a lot over time and you don't have a steady dataset.

1 / 1 point

4. What is a Turing test?

It tests images.

It tests a machine's ability to exhibit intelligent behavior.

It tests the dataset.

It tests and cleans the dataset.

Correct

In 1950, Alan Turing developed the Turing test to test a machine's ability to exhibit intelligent behavior. Alan Turing's test has served as a foundational threshold for artificial intelligence.

1 / 1 point

5. What are some of the different milestones in deep learning history?

Deep Blue defeats a world champion chess player and TensorFlow is released.

Deep Blue defeats a world champion chess player, and Keras is released.

Geoffrey Hinton's work, AlexNet, and TensorFlow

Deep Blue defeats a world champion chess player, and AlexNet is created.

Correct

In 2006, the previous limitations of deep learning, namely exploding and vanishing gradients were overcome with algorithmic advancements such as Geoffrey Hinton's work on unsupervised pre-training. Neural networks are rebranded as deep learning, as we are able to train much deeper networks, networks with more layers. In 2012, a deep learning model using convolutional neural nets called AlexNet achieved a top five error of 15.3 percent. In 2015, one of the most popular libraries, TensorFlow, was built for deep learning, making it more powerful and accessible.

1 / 1 point

6. What is artificial intelligence?

A subset of deep learning.

Any program that can sense, reason, act, and adapt.

A subset of machine learning

None of the above.

Correct

Artificial intelligence is any program that can sense, reason, act, and adapt. It is essentially a machine taking any form of intelligent behavior.

1 / 1 point

7. What are two spaces within AI that are going through drastic growth and innovation?

Computer vision and natural language processing.

Computer vision and deep learning.

Deep learning and machine learning.

Language processing and deep learning.

Correct

In two spaces we are seeing drastic growth and innovation, computer vision and natural language processing. The sharp advancements in computer vision are impacting multiple areas. For example, cars getting to the point where they can drive themselves. Similarly, natural language processing is booming with vast improvements in ability to translate, determine sentiment, clustering of new articles, writing papers, and many others.

1 / 1 point

8. Why did AI flourish so much in the last years?

Faster and inexpensive computers and data storage

Stylish designed computers

Data storage in the cloud is much more expensive

Access to hardware for cleaning data

Correct

AI flourished in the last years with the cloud infrastructure now in place to store large amounts of data for much cheaper, and the plethora of new ways to capture data are now able to build larger, more new once datasets to learn underlying patterns across a multitude of fields. We also have faster computers, and we now have access to powerful hardware for processing and storing data.

1 / 1 point

9. How does Alexa use artificial intelligence?

Recognizes faces and pictures.

Recognizes our voice and answers questions.

Suggests who a person on a photo is.

None of the above answers are correct.

Correct

Alexa, in our homes, recognizes our voice and answers questions or does tasks for us through natural language processing.

1 / 1 point

10. What are the first two steps of a typical machine learning workflow?

Problem statement and data cleaning.

Problem statement and data collection.

Data collection and data transformation.

None of the above answers is correct.

Correct

The first step of a typical machine learning workflow is the problem statement. What problem are you trying to solve? The second step is data collection. What data do you need to solve the problem?

1 / 1 point

https://www.coursera.org/learn/ibm-exploratory-data-analysis-for-machine-learning/assignment-submission/7xNeY/graded-quiz-module-1-modern-ai-a...

1/1